

# Sundaram Singh

 [sundaram26](#) |  [singhsundaram](#) |  [singhsundaram2004@gmail.com](mailto:singhsundaram2004@gmail.com) |  +91-7666713062

## SUMMARY

---

Backend-focused Full-Stack Developer (1.5+ yrs) building production systems for a service marketplace platform — API architecture, multi-gateway payments, and observability — on lean, cost-conscious infrastructure. Currently leading a multi-tenant AI chatbot SaaS from the ground up. Core stack: Node.js/Express, Next.js/React, PostgreSQL, Redis, Qdrant, Docker.

## WORK EXPERIENCE

---

**Full-Stack Developer | Techmasworld Infotech Pvt. Ltd.** May 2025 - Present

- Own backend architecture and API design for a service marketplace platform (Next.js, Node.js, PostgreSQL) serving 25k+ daily requests, within a lean 2–3 person engineering team at an early-stage startup.
- Designed and integrated a multi-gateway payment layer (PhonePe, Razorpay, PayPal) into customizable e-commerce modules, along with Shiprocket logistics integration, analytics dashboards, and admin tooling — built to support merchant-specific configuration without code changes per client.
- Introduced observability into the marketplace platform by standing up Grafana and Prometheus on the production VPS, giving the team visibility into service health and request patterns that didn't exist before.
- Own infrastructure decisions for a cost-constrained, single-VPS deployment (Nginx, Docker) — favoring vertical scaling and operational simplicity over premature distributed infrastructure, appropriate for current traffic and team size.
- Leading development of a multi-tenant AI chatbot SaaS (pre-launch) — architecting knowledge-base integration and embeddable script support for client websites from the ground up.

## PROJECTS

---

**Weave — Multi-Provider AI Agent SDK** [GitHub](#) | [weave.sundaramsingh.com](https://weave.sundaramsingh.com)

- Built a TypeScript AI agent SDK with a runtime state machine, multi-agent handoffs, and a 3-tier memory system (factual, episodic, session) with automated context pruning.
- Engineered an extensible guardrails engine with prompt-injection detection, PII redaction, tool-boundary enforcement, and LLM-based policy classification to control agent/tool execution.
- Designed a provider-agnostic LLM layer supporting OpenAI, Claude, and Gemini with automatic failover and structured-output validation with schema-based self-repair.
- Developed and deployed production documentation and a public landing page, enabling external developers to evaluate and integrate the SDK.

## EDUCATION

---

|                |                                     |                |
|----------------|-------------------------------------|----------------|
| 2030 - present | PhD (Subject) at <b>University</b>  | (GPA: 4.0/4.0) |
| 2023 - 2027    | Bachelor's Degree at <b>College</b> | (GPA: 4.0/4.0) |
| 2022           | Class 12th Some Board               | (Grades)       |
| 2021           | Class 10th Some Board               | (Grades)       |

# PUBLICATIONS

---

# SKILLS

---

|                  |                                                                        |
|------------------|------------------------------------------------------------------------|
| Some Skills      | This, That, Some of this and that etc.                                 |
| Some More Skills | Also some more of this, Some more that, And some of this and that etc. |